

You will use 'Admission\_Predict.csv' for Midterm. This dataset includes the data of the applicants of an academic program. Each application has a unique serial number, which represents a particular student. The dataset contains several parameters which are considered important during the application for Masters Programs. The parameters included are :

- 1) GRE Scores (out of 340)
- 2) TOEFL Scores (out of 120)
- 3) University Rating (out of 5)
- 4) Statement of Purpose (SOP) (out of 5)
- 5) Letter of Recommendation (LOR) Strength (out of 5)
- 6) Undergraduate GPA (out of 10)
- 7) Research Experience (either 0 or 1)
- 8) Chance of Admit (ranging from 0 to 1)

**Q1: Download "Simple linear regression.csv" dataset and load it as 'data'.**

```
In [1]: import pandas as pd
```

```
In [2]: data = pd.read_csv("/content/Simple linear regression.csv")
```

**Q2: Display the first three rows in this dataset.**

```
In [3]: data.head(3)
```

```
Out[3]:
```

	SAT	GPA
0	1714	2.40
1	1664	2.52
2	1760	2.54

```
In [4]: df = data  
df.head(3)
```

```
Out[4]:
```

	SAT	GPA
0	1714	2.40
1	1664	2.52
2	1760	2.54

**Q3: check the duplicate records. if yes then remove the duplicate records.**

```
In [5]: df.duplicated()
```

```
Out[5]: 0      False
        1      False
        2      False
        3      False
        4      False
        ...
        79     False
        80     False
        81     False
        82     False
        83     False
Length: 84, dtype: bool
```

**Q4: Are there any missing values in the dataset?**

```
In [6]: df.isnull().sum()
```

```
Out[6]: SAT      0
        GPA      0
dtype: int64
```

**Q5. Remove the missing value.**

```
In [6]:
```

**Q6: - Display the structure of all variables**

```
In [7]: df.shape
```

```
Out[7]: (84, 2)
```

**Q7: Check the datatypes of the attributes.**

```
In [8]: df.dtypes
```

```
Out[8]: SAT      int64
        GPA      float64
dtype: object
```

## Q8: change the data type according to the discription of the data set

In [9]: *#The data type is according to the distribution of the data set*

## Q9: Print the descriptive statistics of the Simple linear regression data to understand the data a little better (min, max, mean, median, 1st and 3rd quartiles).

In [10]: `df.describe().T`

Out[10]:

	count	mean	std	min	25%	50%	75%	max
<b>SAT</b>	84.0	1845.273810	104.530661	1634.0	1772.00	1846.00	1934.0000	2050.00
<b>GPA</b>	84.0	3.330238	0.271617	2.4	3.19	3.38	3.5025	3.81

## Q10: Divide the dataset to training and test sets.

In [11]: `df = pd.get_dummies(df)  
df.head(2)`

Out[11]:

	SAT	GPA
<b>0</b>	1714	2.40
<b>1</b>	1664	2.52

In [12]: `dfx = df.drop('GPA', axis = 1)`

In [13]: `from sklearn.preprocessing import MinMaxScaler`

In [14]: `scaler = MinMaxScaler()`

```
In [15]: scaler.fit_transform(dfx)
```

```
Out[15]: array([[0.19230769],
                [0.07211538],
                [0.30288462],
                [0.12259615],
                [0.14182692],
                [0.08653846],
                [0.3125    ],
                [0.3125    ],
                [0.37980769],
                [0.51923077],
                [0.24278846],
                [0.33894231],
                [0.24278846],
                [0.1875    ],
                [0.33413462],
                [0.57211538],
                [0.29086538],
                [0.09615385],
                [0.5      ],
                [0.36538462],
                [0.30528846],
                [0.21153846],
                [0.06971154],
                [0.12740385],
                [0.81730769],
                [0.46153846],
                [0.36778846],
                [0.44951923],
                [0.92788462],
                [0.38461538],
                [0.32451923],
                [0.72115385],
                [0.33894231],
                [0.53125   ],
                [0.59134615],
                [0.51682692],
                [0.41826923],
                [0.76923077],
                [0.34375   ],
                [0.47355769],
                [0.55528846],
                [0.51923077],
                [0.79807692],
                [0.16346154],
                [0.85576923],
                [0.69951923],
                [0.45673077],
                [0.77403846],
                [0.53605769],
                [0.82932692],
                [0.40384615],
                [0.53125   ],
                [0.65625   ],
                [0.      ],
                [0.58894231],
                [0.60817308],
                [0.23076923],
```

```
[0.76682692],
[0.35336538],
[0.61778846],
[0.79326923],
[0.41826923],
[0.62259615],
[0.97836538],
[0.62259615],
[0.47596154],
[0.51923077],
[0.72115385],
[0.54567308],
[0.71394231],
[0.71875    ],
[0.34615385],
[0.81971154],
[0.72115385],
[0.93028846],
[0.91586538],
[0.87259615],
[0.92788462],
[0.50240385],
[0.72596154],
[0.42307692],
[0.84855769],
[0.78846154],
[1.         ]])
```

In [16]: `dfx.columns`

Out[16]: `Index(['SAT'], dtype='object')`

In [17]: `pd.DataFrame (scaler.fit_transform(dfx))`

Out[17]:

	0
0	0.192308
1	0.072115
2	0.302885
3	0.122596
4	0.141827
...	...
79	0.725962
80	0.423077
81	0.848558
82	0.788462
83	1.000000

84 rows × 1 columns

```
In [18]: pd.DataFrame(scaler.fit_transform(dfx), columns = dfx.columns)
```

Out[18]:

	SAT
0	0.192308
1	0.072115
2	0.302885
3	0.122596
4	0.141827
...	...
79	0.725962
80	0.423077
81	0.848558
82	0.788462
83	1.000000

84 rows × 1 columns

```
In [19]: scaler_dfx = pd.DataFrame(scaler.fit_transform(dfx), columns = dfx.columns)
scaler_dfx.head(10)
```

Out[19]:

	SAT
0	0.192308
1	0.072115
2	0.302885
3	0.122596
4	0.141827
5	0.086538
6	0.312500
7	0.312500
8	0.379808
9	0.519231

```
In [20]: from sklearn.model_selection import train_test_split
```

```
In [21]: input_train, input_test, output_train, output_test = train_test_split(scaler_dfx,
```

In [22]: input\_train

Out[22]:

	SAT
36	0.418269
27	0.449519
43	0.163462
40	0.555288
2	0.302885
...	...
75	0.915865
9	0.519231
72	0.819712
12	0.242788
37	0.769231

67 rows × 1 columns



```
In [23]: input_train,input_test,output_train,output_test
```

```

Out[23]: (          SAT
          36  0.418269
          27  0.449519
          43  0.163462
          40  0.555288
           2  0.302885
          ..  ...
          75  0.915865
           9  0.519231
          72  0.819712
          12  0.242788
          37  0.769231

[67 rows x 1 columns],
          SAT
          10  0.242788
          67  0.721154
          59  0.617788
          33  0.531250
          83  1.000000
          77  0.927885
          34  0.591346
          38  0.343750
          47  0.774038
          56  0.230769
          78  0.502404
          68  0.545673
          61  0.418269
          66  0.519231
          53  0.000000
          65  0.475962
          31  0.721154,
          SAT
          36  1808
          27  1821
          43  1702
          40  1865
           2  1760
          ..  ...
          75  2015
           9  1850
          72  1975
          12  1735
          37  1954

[67 rows x 1 columns],
          SAT
          10  1735
          67  1934
          59  1891
          33  1855
          83  2050
          77  2020
          34  1880
          38  1777
          47  1956
          56  1730

```

```
78 1843
68 1861
61 1808
66 1850
53 1634
65 1832
31 1934)
```

## Q11 Create the linear regression model.

```
In [24]: import pandas as pd
import numpy as np
```

```
In [25]: from sklearn.linear_model import LinearRegression

model = LinearRegression()

SAT = df[["SAT"]]
GPA = df[["GPA"]]
```

## Q12: Fit the linear regression model on training data set.

```
In [26]: model.fit(SAT, GPA)
```

```
Out[26]: LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

## Q13 : Predicting the test set results

```
In [27]: from sklearn.linear_model import LinearRegression

model = LinearRegression()

SAT = df[["SAT"]]
GPA = df[["GPA"]]

model.fit(SAT, GPA)

X_test = SAT

y_pred = model.predict(X_test)
```

```
In [28]: data_value = 85
```

```
In [29]: predicted = 0.2750402996602799 + 0.00165569 * data_value  
predicted
```

```
Out[29]: 0.4157739496602799
```

### Q14: display regression coefficients

```
In [30]: print(model.intercept_, model.coef_)  
  
0.2750402996602799 [0.00165569]
```